

# diabeties-prediction-eda-svm

November 30, 2023

```
[5]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[6]: df=pd.read_csv('D:/CodeCaulse/diabetes_prediction_dataset.csv')
```

```
[7]: df
```

```
[7]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi
0	Female	80.0	0	1	never	25.19 \
1	Female	54.0	0	0	No Info	27.32
2	Male	28.0	0	0	never	27.32
3	Female	36.0	0	0	current	23.45
4	Male	76.0	1	1	current	20.14
...	...	...	...	...	...	...
99995	Female	80.0	0	0	No Info	27.32
99996	Female	2.0	0	0	No Info	17.37
99997	Male	66.0	0	0	former	27.83
99998	Female	24.0	0	0	never	35.42
99999	Female	57.0	0	0	current	22.43

	HbA1c_level	blood_glucose_level	diabetes
0	6.6	140	0
1	6.6	80	0
2	5.7	158	0
3	5.0	155	0
4	4.8	155	0
...	...	...	...
99995	6.2	90	0
99996	6.5	100	0
99997	5.7	155	0
99998	4.0	100	0
99999	6.6	90	0

[100000 rows x 9 columns]

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null  object
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  object
5   bmi                   100000 non-null  float64
6   HbA1c_level           100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
[9]: print(df.duplicated())
print(df.duplicated().sum())
```

```
0      False
1      False
2      False
3      False
4      False
...
99995   True
99996  False
99997  False
99998  False
99999  False
Length: 100000, dtype: bool
3854
```

```
[10]: df.drop_duplicates(inplace=True)
```

```
[11]: df.duplicated().sum()
```

```
[11]: 0
```

```
[12]: df.isna().sum()
```

```
[12]: gender          0
      age            0
      hypertension    0
      heart_disease    0
      smoking_history  0
      bmi            0
      HbA1c_level     0
      blood_glucose_level 0
      diabetes        0
      dtype: int64
```

```
[13]: df.describe(include='object')
```

```
[13]:      gender smoking_history
count    96146          96146
unique         3           6
top    Female          never
freq     56161          34398
```

```
[14]: df.describe()
```

```
[14]:      age hypertension heart_disease      bmi  HbA1c_level \
count  96146.000000  96146.000000  96146.000000  96146.000000  96146.000000 \
mean    41.794326    0.077601    0.040803    27.321461    5.532609
std     22.462948    0.267544    0.197833     6.767716    1.073232
min      0.080000    0.000000    0.000000    10.010000    3.500000
25%     24.000000    0.000000    0.000000    23.400000    4.800000
50%     43.000000    0.000000    0.000000    27.320000    5.800000
75%     59.000000    0.000000    0.000000    29.860000    6.200000
max     80.000000    1.000000    1.000000    95.690000    9.000000

      blood_glucose_level  diabetes
count    96146.000000  96146.000000
mean     138.218231    0.088220
std       40.909771    0.283616
min       80.000000    0.000000
25%      100.000000    0.000000
50%      140.000000    0.000000
75%      159.000000    0.000000
max      300.000000    1.000000
```

```
[15]: numeric_col=[]
      non_numeric_col=[]
      for column in df.columns:
          if pd.api.types.is_numeric_dtype(df[column]):
              if(df[column].nunique()<5):
                  non_numeric_col.append(column)
```

```

        else:
            numeric_col.append(column)
    else:
        non_numeric_col.append(column)

```

```

[16]: print(numeric_col)
      print(non_numeric_col)

```

```

['age', 'bmi', 'HbA1c_level', 'blood_glucose_level']
['gender', 'hypertension', 'heart_disease', 'smoking_history', 'diabetes']

```

```

[17]: def univariate_analysis_numeric(col):
      fig, ax = plt.subplots(1, 2, figsize=(12,10))
      sns.histplot(x=df[col], kde=True, bins=20, color='skyblue', ax=ax[0])
      ax[0].set_title(f'Histogram of {col}.')

      sns.boxplot(x=df[col], ax=ax[1])
      ax[1].set_title(f'Boxplot diagram of {col}      ')

```

```



[18]: for col in numeric_col:
      print(f' Univariate analysis for {col} column:')
      univariate_analysis_numeric(col)



```



```


Univariate analysis for age column:
Univariate analysis for bmi column:
Univariate analysis for HbA1c_level column:
Univariate analysis for blood_glucose_level column:

```









```
[19]: def univariate_analysis_cat(col):
    fig, ax = plt.subplots(1, 2, figsize=(12, 5))

    # Countplot
    sns.countplot(x=df[col], data=df, palette='viridis', ax=ax[0])
    ax[0].set_title(f'Countplot for {col}')



    # Pie plot
    data_counts = df[col].value_counts()
    ax[1].pie(data_counts, labels=data_counts.index, autopct='%1.3f%%',
    ↪ startangle=90, colors=sns.color_palette('pastel'))
    ax[1].set_title(f'Pie plot for {col}')

    plt.show()



[20]: for col in non_numeric_col:
    print(f' Univariate analysis for {col} column:')
    univariate_analysis_cat(col)
```





Univariate analysis for gender column:




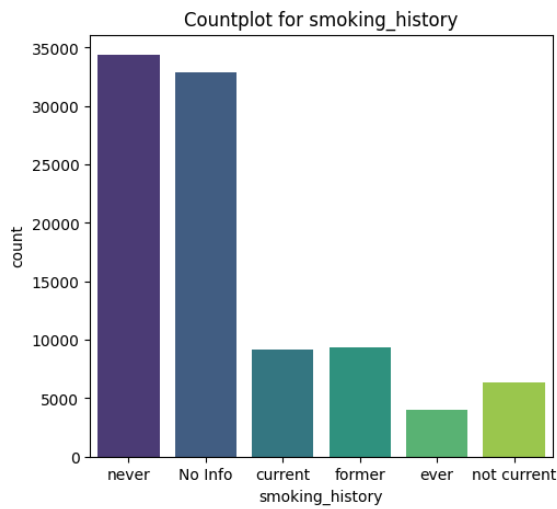
Univariate analysis for hypertension column:





Univariate analysis for heart\_disease column:



Univariate analysis for smoking\_history column:



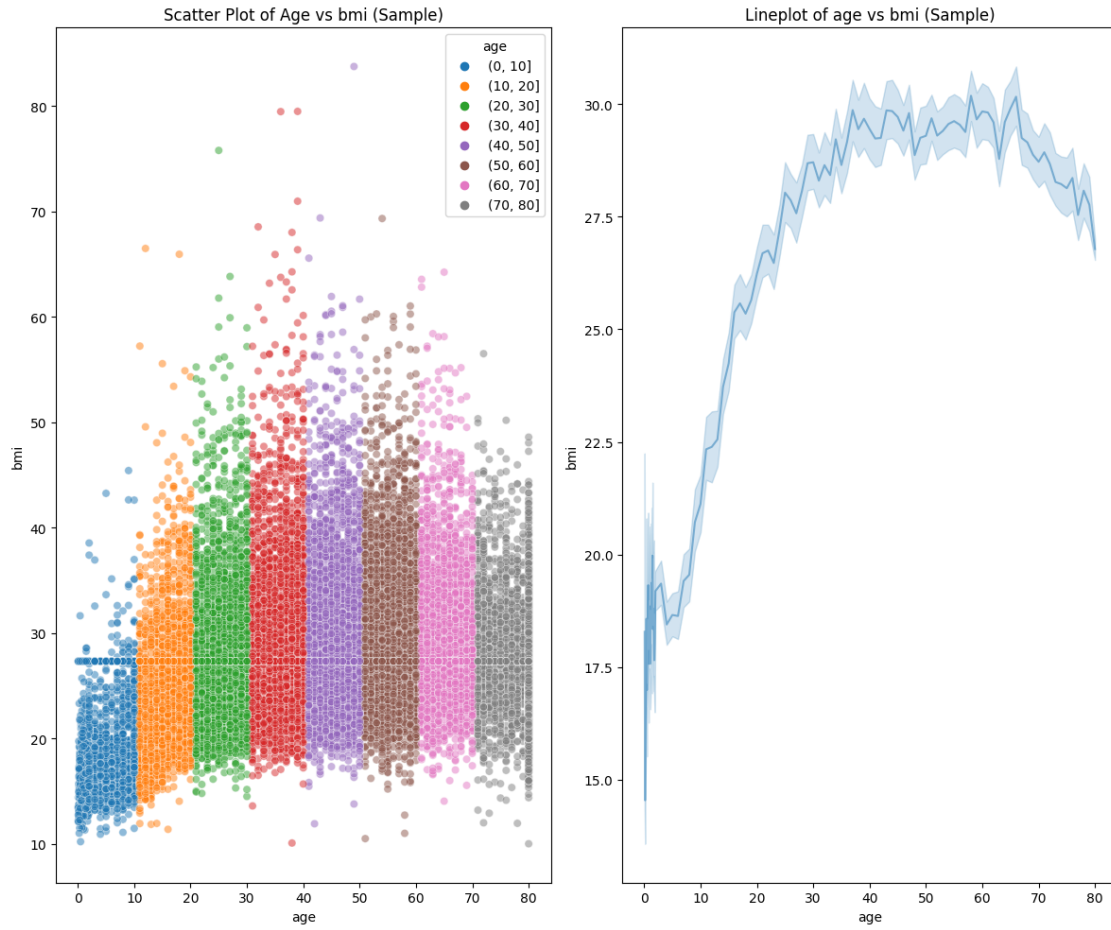
Univariate analysis for diabetes column:

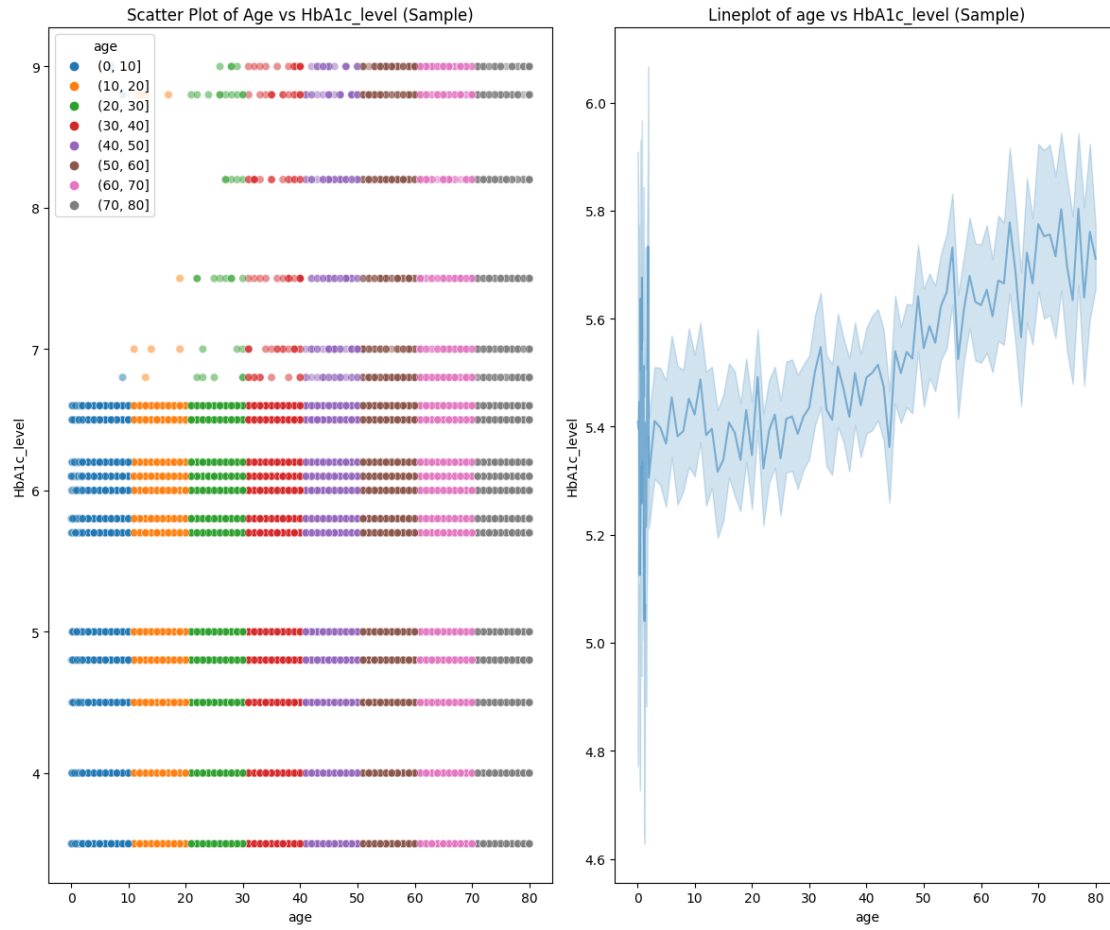



```
[21]: sample_size = 30000
sample_df = df.sample(n=sample_size, random_state=42)
age_group = pd.cut(df['age'], bins=[0, 10, 20, 30, 40, 50, 60, 70, 80])

for col in numeric_col[1:]:
    fig, ax = plt.subplots(1, 2, figsize=(12, 10))
    sns.scatterplot(x='age', y=df[col], data=sample_df, hue=age_group, alpha=0.
↪5, ax=ax[0])
    ax[0].set_title(f'Scatter Plot of Age vs {col} (Sample)')

    sns.lineplot(x='age', y=df[col], data=sample_df, alpha=0.5, ax=ax[1])
    ax[1].set_title(f'Lineplot of age vs {col} (Sample)')
    plt.tight_layout()
    plt.show()
```









```
[22]: print(sample_df['age'].mean())
      print(df['age'].mean())
```

```
41.856903999999999
```


```
41.79432571297817
```


```
[23]: age_group = pd.cut(df['age'], bins=[0, 10, 20, 30, 40, 50, 60, 70, 80])
      for i in non_numeric_col:
          plt.figure(figsize=(12,10))
          sns.countplot(x=df[i],data=df,hue=age_group)
          plt.title(f'Age distribution vs {i}')
          plt.show()
```














```
[24]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for col in non_numeric_col:
    sample_df[col]=le.fit_transform(sample_df[col])
sample_df
```

```
[24]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi
2547	0	5.0	0	0	0	15.14 \
34774	1	19.0	0	0	4	22.74
71084	0	52.0	0	0	4	24.27
50584	1	55.0	0	0	0	27.32
80788	0	22.0	0	0	4	28.78
...	...	...	...	...	...	...
54431	0	26.0	0	0	3	28.27
55182	0	57.0	0	0	4	22.12
92776	1	74.0	0	1	2	29.83
76035	0	11.0	0	0	0	27.32
81939	0	80.0	0	0	1	20.03

	HbA1c_level	blood_glucose_level	diabetes
2547	4.5	90	0
34774	6.6	85	0
71084	7.5	220	1
50584	5.7	85	0
80788	4.5	200	0
...	...	...	...
54431	6.2	100	0
55182	5.0	100	0
92776	4.8	160	0
76035	6.2	90	0
81939	6.6	90	0

[30000 rows x 9 columns]

```
[25]: x = sample_df.drop('diabetes', axis=1)
      y = sample_df['diabetes']
```

```
[26]: from sklearn.svm import SVC
      from sklearn.metrics import classification_report, accuracy_score,
      ↪confusion_matrix
      X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.3)
      model_svm = SVC(kernel = 'linear', random_state = 0)
      model_svm.fit(X_train, y_train)
```

```
[26]: SVC(kernel='linear', random_state=0)
```

```
[27]: y_pred = model_svm.predict(X_test)
```

```
[28]: accuracy_svm = accuracy_score(y_test, y_pred)
      print(accuracy_svm)
```

0.9582222222222222

```
[29]: classification_rep_svm = classification_report(y_test, y_pred)
      print(classification_rep_svm)
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	8192
1	0.93	0.58	0.71	808
accuracy			0.96	9000
macro avg	0.95	0.79	0.85	9000
weighted avg	0.96	0.96	0.95	9000